


# Caching long lasting queries

Laszlo Forro @ thinkproject.com



PostgreSQL

**thinkproject**

#constructionintelligence


---

Serve *It Fresh* and if we are not *Fast* fast enough – Cache *It*

Rules of caching on the back end:

- 1) Do not cache*
- 2) Cache is last resort, so try not to cache*
- 3) Before caching consider if there is better data model or architecture*

## Side note



- the database and the user are separated
- the entities are how the application developers imagine it
- cognitive gap

Gavin Powell: Beginning Database Design, Wiley, 2006,  
Ch.2.: Database Modelling in the Workspace

---

## Serve *It* Fresh and Fast

*It*: what is the *use pattern* of the data we need to serve?

- Can we say that the most recently used?
- Can we say that the most frequently used?
- Or neither of those above?

---

## Serve *It* Fresh and Fast

*It* : what is the *identity* of the data we serve *regarding the data source*?

- Can be identified simply
- Or a product of function like any aggregate (sum, count, aggregate average ), max, min
- Or can't be identified, eg. a result of a DISTINCT

## Serve It Fresh and Fast

Fresh : what is the allowed *difference* of the data we serve and the *data source*?

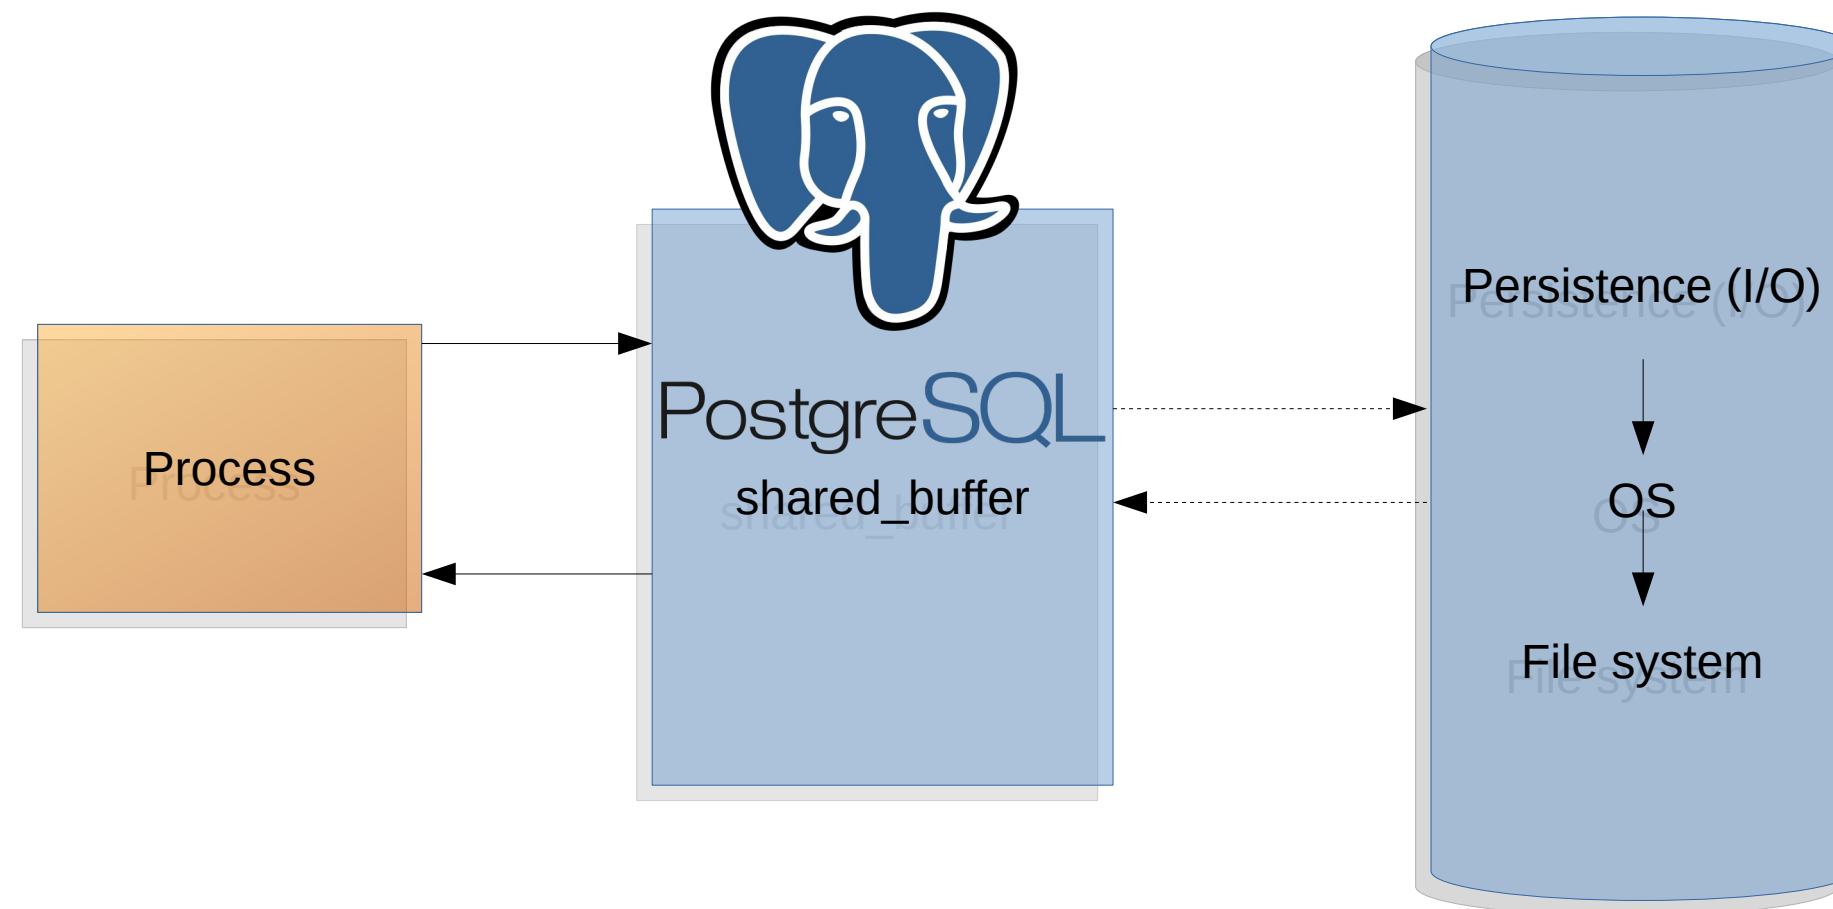
- Do we have to serve the latest and greatest for everyone
- Or we can let the user trigger an action for himself?
- What latency can we permit?

## Serve It Fresh and *Fast*

*Fast* : which *system component* is identified as 'slow'?

- Where should the optimization take place?
- Can a change in the architecture/data model result in better optimization ?
- What is the implication of Cache Invalidation requirements ?  
(change ratio, cache complexity)

## Serve It From Cache



- Caching in shared\_buffer
- Clock-sweep using access counter to release memory
- Managing inserts and updates
- Low level on blocks / tuples
- Most Frequently Used/  
Most Recently Used
- It is about *all* access to the database, not only that we consider user relevant.

If we need cache, it is already not sufficient...

See the excellent chapter on [PostgreSQL buffer management](#).

See also PostgreSQL memory components @severalnines: [Architecture and Tuning of Memory in PostgreSQL Databases](#)

See also the [Cache replacement policies](#) page on the Wikipedia.



# Serve *It* From *Cache*

## *Incremental View Maintenance*

A materialized views represent a result set of a query.

On change try to add only the change instead of replaying the query.



Assumption:

- the change can be propagated, because there are entities that can be identified as changed
- the change propagation is cheaper than replaying the query
  - factors: change ratio, algorithmic efficiency of view delegation

Serve *It* From *Cache*  
*Incremental View Maintenance*

A *very* simple example:

1. Let's say a view is  $V = A \bowtie_{\theta} B$
2. Change in A means a V diff table:  $V_{\delta} = A_{\delta} \bowtie_{\theta} B$
3. Add the changes of  $V_{\delta}$  to V using some Update Propagation Query




Serve *It* From *Cache*  
*Incremental View Maintenance*

UPQ for INSERT can be  $V \cup V_\delta$  or  $V - V_\delta$  on DELETE,  
 $(V - V_\delta) \cup V_\delta$  on UPDATE *generally speaking*.



Serve *It* From *Cache*  
*IVM – Self Maintainable View*

- there is a relation in the View and Update/Delete definition that makes possible to Update/Delete the View without replaying the whole query
- Inserts must be still evaluated



Serve *It* From *Cache*  
*IVM – Self Maintainable View*


The *very* simple example

$$V = A \bowtie_{\theta} B \rightarrow V_{\delta} = A_{\delta} \bowtie_{\theta} B \rightarrow \text{Update } V \text{ with } V_{\delta}$$

would look like

1.  $V = A \bowtie_{a.id = b.aid} B$
2. Then when updating  $A_{a1..aN}$ ,  $V$  can also be updated on the same ID references

Eg.: Katsis et al. : Utilizing IDs to Accelerate Incremental View Maintenance




# Serve *It* From *Cache*

## *IVM – Self Maintainable View*

A possible algorithm for PostgreSQL:

- on change of an entity in table A check
  1. if there is a dependent MV view on A and if yes,
  2. check if that view has a column corresponding to A.id
- make an UPDATE ... RETURNING A.id on A
- UPDATE the view with the NEW values by A.id
- might still need to parse a good amount of SQL
  - eg. A.id = B.aid (foreign key)
  - SELECT A.id as foo

Etc.



Serve *It* From *Cache*  
*IVM – Self Maintainable  
View*

Yugo Nagata: Incremental View Maintenance,  
presented at 2018 PgConf EU, Lissabon.  
Implemented as an [experimental PostgreSQL fork](#).



- The problem to identify entities is solved by using the OIDs
- Use a mapping table between the OIDs of the source table tuples and the View tuples for each tuple

# Serve It From Cache

## *Incremental View Maintenance*

- for a generic solution we need to identify an Update Propagation Query
- it grows more and more complex as more and more tables are involved
- needs efficient algorithm for DISTINCT, aggregate, min-max etc. functions
- how can it be efficiently implemented?



*Database Query Optimization problem!*

For the generic theory and the query optimizations see the thesis of Dimitra Vista:

[Optimizing Incremental View Maintenance Expressions in Relational Databases, Toronto, 1996.](#)

**thinkproject**

#constructionintelligence




# Serve *It* From *Cache*

## *Incremental View Maintenance*

The approaches above

- create a Deltas table:
  - that needs maintenance
- need some IDs to identify the changed entities:
  - that is not always available for the whole View Definition
- apply the changes via some automated method
- track parallel changes...

Leads to a complex query optimization problem.



# Serve *It* From *Cache*

## *A Logical Decoding Based Solution*

- listener on logical decoding socket filtering the stream
- change distributor forwards the data to workers (domain filtering and semantics)
- the worker output is the cached data structure

Blagoj Atanasovski: [Use Logical Decoding to build your own application cache](#), Presented at PgConf EU 2018, Lissabon

**thinkproject**

#constructionintelligence

# Serve *It* From *Cache*

## *A Logical Decoding Based Solution*

### Advantages:

- Consistency and invalidation is trivial
- No complex update query problem
- Application developer friendly

Blagoj Atanasovski: [Use Logical Decoding to build your own application cache](#), Presented at PgConf EU 2018, Lissabon

**thinkproject**

#constructionintelligence

## Serve *It* From *Cache*

*Our scenario:*

- dynamically generated complex queries
- the identity of the entities depend on the query and not DB relations
- the use case does not define hot data
- the only correlation is: the bigger a result set the more INSERTS
- no DELETE
- must be always *Fresh*
- max 1 sec user wait is ok.

## Serve It From Cache

### *Incremental View Maintenance*

A different angle: a change TS is always incremental.  
Let A have an attribute,  $t$ , that acts as a 'last\_changed' TS.


$T_0$ : Let's say a view is  $V = A \bowtie_{\theta} B$

Save  $T_0$  as metadata for the View update TS.

$T_1$ : Update A as usual, but the  $t = T_1$  for each element of  $A_{\delta}$

$T_1$ : Update V:  $V \leftarrow (\sigma_{(t > T_0)} A) \bowtie_{\theta} B$ , using  $T_0$  from the View metadata table.

Save  $T_1$  as metadata for the View update TS.



# Serve It From Cache

## *Incremental View Maintenance*


To identify an entity to Update, a Unique key is needed on the V table.

INSERT – UPDATE implementation: UPSERT.

If there is a key violation, it means that the entity already exists in the V table,  
so it needs UPDATE.

What the client must provide:

- some way to recognize the condition for TS check in the V query
- the Unique key for the V table




## Serve It From Cache

query\_cache : *non self maintaining view with deferred update*

What we have not solved / did not want to solve  
/ avoided to solve / couldn't solve / did not even think of to solve:

- a generic solution for the Update Propagation Query, instead we let the Query Optimizer of the Database to do the job.
- heruistics: the user must provide the Unique key.
- DELETION: it must be implemented as a change.
- optimization for aggregates
- ...




Serve *It* From *Cache*

*query\_cache* : *non self maintaining view with deferred update*


TODOs:

- implement DELETE
- implement AGGREGATE, MIN/MAX, COUNT, AVG etc. cache
- rethink of the API
- some automatism (eg. if outer query has DISTINCT that can act as Unique key)
- possibly make it more generic





Thank you !



PostgreSQL

**thinkproject**